

```
import pandas as pd
import numpy as np
from sklearn.preprocessing import LabelEncoder
```

```
liver_df= pd.read_csv('/content/indian_liver_patient.csv')
liver_df.head()
```



	Age	Gender	Total_Bilirubin	Direct_Bilirubin	Alkaline_Phosphotase	Alamine_Aminot
0	65	Female	0.7	0.1	187	
1	62	Male	10.9	5.5	699	
2	62	Male	7.3	4.1	490	
3	58	Male	1.0	0.4	182	
4	72	Male	3.9	2.0	195	

Next steps:

[Generate code with liver_df](#)[View recommended plots](#)

```
liver_df.head()
```



	Age	Gender	Total_Bilirubin	Direct_Bilirubin	Alkaline_Phosphotase	Alamine_Aminot
0	65	Female	0.7	0.1	187	
1	62	Male	10.9	5.5	699	
2	62	Male	7.3	4.1	490	
3	58	Male	1.0	0.4	182	
4	72	Male	3.9	2.0	195	

Next steps:

[Generate code with liver_df](#)[View recommended plots](#)

```
liver_df.info()
```



```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 583 entries, 0 to 582
Data columns (total 11 columns):
#   Column              Non-Null Count  Dtype
---  -
0   Age                  583 non-null   int64
1   Gender               583 non-null   object
```

```

2   Total_Bilirubin      583 non-null    float64
3   Direct_Bilirubin     583 non-null    float64
4   Alkaline_Phosphotase 583 non-null    int64
5   Alamine_Aminotransferase 583 non-null    int64
6   Aspartate_Aminotransferase 583 non-null    int64
7   Total_Protiens       583 non-null    float64
8   Albumin              583 non-null    float64
9   Albumin_and_Globulin_Ratio 579 non-null    float64
10  Dataset              583 non-null    int64

```

dtypes: float64(5), int64(5), object(1)

memory usage: 50.2+ KB

```
liver_df.describe(include='all')
```



	Age	Gender	Total_Bilirubin	Direct_Bilirubin	Alkaline_Phosphotase	Ala
count	583.000000	583	583.000000	583.000000	583.000000	
unique	NaN	2	NaN	NaN	NaN	
top	NaN	Male	NaN	NaN	NaN	
freq	NaN	441	NaN	NaN	NaN	
mean	44.746141	NaN	3.298799	1.486106	290.576329	
std	16.189833	NaN	6.209522	2.808498	242.937989	
min	4.000000	NaN	0.400000	0.100000	63.000000	
25%	33.000000	NaN	0.800000	0.200000	175.500000	
50%	45.000000	NaN	1.000000	0.300000	208.000000	
75%	58.000000	NaN	2.600000	1.300000	298.000000	
max	90.000000	NaN	75.000000	19.700000	2110.000000	

```
liver_df.columns
```



```

Index(['Age', 'Gender', 'Total_Bilirubin', 'Direct_Bilirubin',
      'Alkaline_Phosphotase', 'Alamine_Aminotransferase',
      'Aspartate_Aminotransferase', 'Total_Protiens', 'Albumin',
      'Albumin_and_Globulin_Ratio', 'Dataset'],
      dtype='object')

```

```
liver_df.isnull().sum()
```



```

Age      0
Gender    0
Total_Bilirubin    0
Direct_Bilirubin    0
Alkaline_Phosphotase    0

```

```

Alamine_Aminotransferase    0
Aspartate_Aminotransferase  0
Total_Protiens              0
Albumin                     0
Albumin_and_Globulin_Ratio  4
Dataset                     0
dtype: int64

```

```
pd.get_dummies(liver_df['Gender'], prefix = 'Gender').head()
```



	Gender_Female	Gender_Male
0	True	False
1	False	True
2	False	True
3	False	True
4	False	True



```
liver_df = pd.concat([liver_df,pd.get_dummies(liver_df['Gender'], prefix = 'Gender')], axis
```

```
liver_df.head()
```



	Age	Gender	Total_Bilirubin	Direct_Bilirubin	Alkaline_Phosphotase	Alamine_Aminot
0	65	Female	0.7	0.1	187	
1	62	Male	10.9	5.5	699	
2	62	Male	7.3	4.1	490	
3	58	Male	1.0	0.4	182	
4	72	Male	3.9	2.0	195	

Next steps: [Generate code with liver_df](#)

[View recommended plots](#)

```
liver_df.describe()
```



	Age	Total_Bilirubin	Direct_Bilirubin	Alkaline_Phosphotase	Alamine_Ami
count	583.000000	583.000000	583.000000	583.000000	
mean	44.746141	3.298799	1.486106	290.576329	
std	16.189833	6.209522	2.808498	242.937989	
min	4.000000	0.400000	0.100000	63.000000	
25%	33.000000	0.800000	0.200000	175.500000	
50%	45.000000	1.000000	0.300000	208.000000	
75%	58.000000	2.600000	1.300000	298.000000	
max	90.000000	75.000000	19.700000	2110.000000	

```
liver_df[liver_df['Albumin_and_Globulin_Ratio'].isnull()]
```



	Age	Gender	Total_Bilirubin	Direct_Bilirubin	Alkaline_Phosphotase	Alamine_Amin
209	45	Female	0.9	0.3	189	
241	51	Male	0.8	0.2	230	
253	35	Female	0.6	0.2	180	
312	27	Male	1.3	0.6	106	

```
liver_df["Albumin_and_Globulin_Ratio"] = liver_df.Albumin_and_Globulin_Ratio.fillna(liver_d
```

```
X = liver_df.drop(['Gender', 'Dataset'], axis=1)
X.head(3)
```



	Age	Total_Bilirubin	Direct_Bilirubin	Alkaline_Phosphotase	Alamine_Aminotransfera
0	65	0.7	0.1	187	
1	62	10.9	5.5	699	
2	62	7.3	4.1	490	

Next steps:

[Generate code with X](#)

[View recommended plots](#)

```
y = liver_df['Dataset']
```

```
liver_corr = X.corr()
liver_corr
```



	Age	Total_Bilirubin	Direct_Bilirubin	Alkaline_Phosp
Age	1.000000	0.011763	0.007529	0.080425
Total_Bilirubin	0.011763	1.000000	0.874618	0.206669
Direct_Bilirubin	0.007529	0.874618	1.000000	0.234939
Alkaline_Phosphatase	0.080425	0.206669	0.234939	1.000000
Alamine_Aminotransferase	-0.086883	0.214065	0.233894	0.086883
Aspartate_Aminotransferase	-0.019910	0.237831	0.257544	0.019910
Total_Protiens	-0.187461	-0.008099	-0.000139	-0.187461
Albumin	-0.265924	-0.222250	-0.228531	-0.265924
Albumin_and_Globulin_Ratio	-0.216089	-0.206159	-0.200004	-0.216089
Gender_Female	-0.056560	-0.089291	-0.100436	0.056560
Gender_Male	0.056560	0.089291	0.100436	-0.056560

Next steps:

[Generate code with liver_corr](#)
[View recommended plots](#)

```
from sklearn.metrics import accuracy_score
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report, confusion_matrix
from sklearn import linear_model
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC, LinearSVC
from sklearn.ensemble import RandomForestClassifier, AdaBoostClassifier, BaggingClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.linear_model import Perceptron
from sklearn.linear_model import SGDClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.neural_network import MLPClassifier
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.30, random_state=101)
print (X_train.shape)
print (y_train.shape)
print (X_test.shape)
print (y_test.shape)
```



```
(408, 11)
(408,)
(175, 11)
```

(175,)

Logistic Regression

```
logreg = LogisticRegression()
logreg.fit(X_train, y_train)
```

→ /usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py:458: Converge
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

```
    ▾ LogisticRegression
```

```
    LogisticRegression())
```

```
log_predicted= logreg.predict(X_test)
logreg_score = round(logreg.score(X_train, y_train) * 100, 2)
logreg_score_test = round(logreg.score(X_test, y_test) * 100, 2)
```

```
print('Logistic Regression Training Score: \n', logreg_score)
print('Logistic Regression Test Score: \n', logreg_score_test)
print('Coefficient: \n', logreg.coef_)
print('Intercept: \n', logreg.intercept_)
print('Accuracy: \n', accuracy_score(y_test, log_predicted))
print('Confusion Matrix: \n', confusion_matrix(y_test, log_predicted))
print('Classification Report: \n', classification_report(y_test, log_predicted))
```

→ Logistic Regression Training Score:
72.06
Logistic Regression Test Score:
68.0
Coefficient:
[[-0.00994992 -0.0985122 -0.30688724 -0.00082939 -0.01078827 -0.00275598
-0.23899684 0.40208926 0.59475501 0.2533529 0.0911599]]
Intercept:
[0.36100669]
Accuracy:
0.68
Confusion Matrix:
[[107 17]
 [39 12]]
Classification Report:


	precision	recall	f1-score	support
1	0.73	0.86	0.79	124
2	0.41	0.24	0.30	51

accuracy			0.68	175
macro avg	0.57	0.55	0.55	175
weighted avg	0.64	0.68	0.65	175

```
coeff_df = pd.DataFrame(X.columns)
coeff_df.columns = ['Feature']
coeff_df["Correlation"] = pd.Series(logreg.coef_[0])
pd.Series(logreg.coef_[0])

coeff_df.sort_values(by='Correlation', ascending=False)
```



	Feature	Correlation	
8	Albumin_and_Globulin_Ratio	0.594755	
7	Albumin	0.402089	
9	Gender_Female	0.253353	
10	Gender_Male	0.091160	
3	Alkaline_Phosphotase	-0.000829	
5	Aspartate_Aminotransferase	-0.002756	
0	Age	-0.009950	
4	Alamine_Aminotransferase	-0.010788	
1	Total_Bilirubin	-0.098512	
6	Total_Protiens	-0.238997	
2	Direct_Bilirubin	-0.306887	

Gaussian Naive Bayes

```
gaussian = GaussianNB()
gaussian.fit(X_train, y_train)
gauss_predicted = gaussian.predict(X_test)
```

```
gauss_score = round(gaussian.score(X_train, y_train) * 100, 2)
gauss_score_test = round(gaussian.score(X_test, y_test) * 100, 2)
print('Gaussian Score: \n', gauss_score)
print('Gaussian Test Score: \n', gauss_score_test)
print('Accuracy: \n', accuracy_score(y_test, gauss_predicted))
print(confusion_matrix(y_test, gauss_predicted))
print(classification_report(y_test, gauss_predicted))
```



Gaussian Score:

56.13

Gaussian Test Score:

53.14

Accuracy:

0.5314285714285715

[[44 80]

[2 49]]

	precision	recall	f1-score	support
1	0.96	0.35	0.52	124
2	0.38	0.96	0.54	51
accuracy			0.53	175
macro avg	0.67	0.66	0.53	175
weighted avg	0.79	0.53	0.53	175

Random Forest

```
random_forest = RandomForestClassifier(max_depth=3,n_estimators=56,criterion='entropy')
random_forest.fit(X_train, y_train)
```



RandomForestClassifier

```
RandomForestClassifier(criterion='entropy', max_depth=3, n_estimators=56)
```

```
rf_predicted = random_forest.predict(X_test)
```

```
random_forest_score = round(random_forest.score(X_train, y_train) * 100, 2)
random_forest_score_test = round(random_forest.score(X_test, y_test) * 100, 2)
print('Random Forest Score: \n', random_forest_score)
print('Random Forest Test Score: \n', random_forest_score_test)
print('Accuracy: \n', accuracy_score(y_test,rf_predicted))
print(confusion_matrix(y_test,rf_predicted))
print(classification_report(y_test,rf_predicted))
```



Random Forest Score:

76.23

Random Forest Test Score:

70.29

Accuracy:

0.7028571428571428


[[117 7]

[45 6]]



	precision	recall	f1-score	support
1	0.72	0.94	0.82	124
2	0.46	0.12	0.19	51
accuracy			0.70	175

macro avg	0.59	0.53	0.50	175
weighted avg	0.65	0.70	0.63	175

```
finX = liver_df[['Total_Protiens', 'Albumin', 'Gender_Male']]
finX.head(4)
```



	Total_Protiens	Albumin	Gender_Male
0	6.8	3.3	False
1	7.5	3.2	True
2	7.0	3.3	True
3	6.8	3.4	True

Next steps:


[Generate code with finX](#)[View recommended plots](#)

Logistic Regression

```
X_train, X_test, y_train, y_test = train_test_split(finX, y, test_size=0.30, random_state=1
```

```
logreg = LogisticRegression()
```

```
logreg.fit(X_train, y_train)
```



▾ LogisticRegression
 LogisticRegression()

```
log_predicted= logreg.predict(X_test)
```

```
logreg_score = round(logreg.score(X_train, y_train) * 100, 2)
logreg_score_test = round(logreg.score(X_test, y_test) * 100, 2)
```

```
# Equation coefficient and Intercept
```

```
print('Logistic Regression Training Score: \n', logreg_score)
print('Logistic Regression Test Score: \n', logreg_score_test)
print('Coefficient: \n', logreg.coef_)
print('Intercept: \n', logreg.intercept_)
print('Accuracy: \n', accuracy_score(y_test, log_predicted))
print('Confusion Matrix: \n', confusion_matrix(y_test, log_predicted))
print('Classification Report: \n', classification_report(y_test, log_predicted))
```

```

Logistic Regression Training Score:
71.08
Logistic Regression Test Score:
71.43
Coefficient:
[[-0.58254377  1.08940697 -0.54185122]]
Intercept:
[-0.20423275]
Accuracy:
0.7142857142857143
Confusion Matrix:
[[120   4]
 [ 46   5]]
Classification Report:

```

	precision	recall	f1-score	support
1	0.72	0.97	0.83	124
2	0.56	0.10	0.17	51
accuracy			0.71	175
macro avg	0.64	0.53	0.50	175
weighted avg	0.67	0.71	0.63	175

Decision Tree Classifier

```
dt=DecisionTreeClassifier()
```

```
dt.fit(X_train,y_train)
```

```

DecisionTreeClassifier
DecisionTreeClassifier()

```

```
y_pred=dt.predict(X_test)
```

```

dt_score = round(dt.score(X_train, y_train) * 100, 2)
dt_test = round(dt.score(X_test, y_test) * 100, 2)

```

```

from sklearn.metrics import accuracy_score
accuracy_score(y_test,y_pred)

```

```
0.6457142857142857
```

```

from sklearn.metrics import confusion_matrix
confusion_matrix(y_test,y_pred)

```

```

array([[100, 24],
       [ 38, 13]])

```

Model evaluation

```
models = pd.DataFrame({
    'Model': [ 'Logistic Regression', 'Gaussian Naive Bayes', 'Random Forest', 'Decision Tree']
    'Score': [ logreg_score, gauss_score, random_forest_score, dt_score],
    'Test Score': [ logreg_score_test, gauss_score_test, random_forest_score_test, dt_test]})
models.sort_values(by='Test Score', ascending=False)
```



	Model	Score	Test Score	
0	Logistic Regression	71.08	71.43	
2	Random Forest	76.23	70.29	
3	Decision Tree	93.38	64.57	
1	Gaussian Naive Bayes	56.13	53.14	



```
import pickle
```

```
filename = 'liver.sav'
pickle.dump(round, open(filename, 'wb'))
```

```
# loading the saved model
loaded_model = pickle.load(open('liver.sav', 'rb'))
```

Start coding or [generate](#) with AI.